

Changes made to the evaluation methodology, the models tested, and the reported figures, following the reviewers' report

Change	Before	After	Reason
EVALUATION SET — THE ANSWER TO THE CO-DESIGN-BIAS CRITICISM			
Attack generator	8 tiers × 10 hand-written strings = 80 unique attacks, sampled with replacement. One string appeared 25 times in a single 1,001-attack batch.	Combinatorial <i>template × variable-pool</i> architecture: 1,455 unique base strings, sampled without replacement. Zero duplicates across all five seeds. Over-requesting a tier now raises an error instead of repeating.	A 1,001-attack run drawn from 80 strings is a test of 80 attacks replicated twelve times: the effective sample size was 80, so every confidence interval computed on n=1,001 was invalid, and detection could not be distinguished from memorising the author's own strings.
Benign query generator	12 hand-written questions, sampled with replacement (~28 copies of each per 333-query batch). None contained a word the detector treats as risky.	Combinatorial pools, 333 unique queries per batch, deliberately including benign-but-sensitive phrasings (e.g. "how do I <i>override</i> a method", ".gitignore configuration"). Asserted by an automated check.	The old set was structurally incapable of producing a false positive , so the reported FPR of 0.00 % was not a measurement of anything. This is the single most important correction in the revision.
L2 rule patterns	Patterns fitted, unintentionally, to the 80 written strings. On a diverse batch pre-L4 detection fell to 60.1 % .	Patterns re-derived from each attack tier's own enumerable structure — never from held-out data. Pre-L4 detection 88.0 % .	The 8.4-point drop from 68.5 % to 60.1 % when the generator was replaced is the measured size of the co-design bias , observed inside the project's own evaluation. Coverage was then rebuilt on structure rather than vocabulary.
METRICS AND DEFINITIONS			
Attack success rate definition	ASR was computed as the complement of the detection rate. Every response that reached the model was counted as a successful attack.	Reaching the model and complying are now measured separately. External: 57.40 % reached the model, 15.92 % actually complied (Mistral-7B). 70.3 % of attacks that passed every defense layer were refused by the model itself.	The second reviewer criticism. Equating reach with success overstates the attack success rate by a factor of 3.6× . Compliance is judged by comparing the response against the literal injected instruction and against the real question, with ambiguous cases counted as <i>not</i> complied — so the figure is a lower bound.
Injected-instruction recovery	A line-split heuristic that only worked for <i>start</i> and <i>end</i> positions — roughly 70 % coverage, with mid-document injections silently dropped.	Ground-truth lookup of each sample's attack identifier against BIPiA's own source files. 100 % coverage of reached-model rows for both models.	The excluded rows were being counted as non-compliant by default, and the per-verdict percentages divided by the smaller extractable count. The literal source text replaces the heuristic entirely.
Response text used for scoring	A 500-character snippet. 52.4 % of responses were truncated.	The full untruncated response.	Truncation depressed the benign-similarity term (0.297 vs 0.556) and therefore over-reported compliance . Re-measuring on full text was a prerequisite for trusting the number.
Layer contribution table	Derived indirectly from cumulative ablation deltas: L1 195 / L2 640 / L3 34 / L4 41.	Direct per-query tally of which layer actually fired: L1 158 / L2 720 / L3 24 / L4 10 (912 blocked + 89 reached = 1,001). Both methods now agree on all four layers.	A derived figure was being reported as a measurement. The direct tally is the primary number; the ablation deltas are retained as an independent cross-check, and their exact agreement validates both.
MODELS, RUNS AND REPRODUCIBILITY			
Models evaluated	One — Mistral-7B-Instruct v0.2.	Two — Mistral-7B-Instruct v0.2 (primary) and Llama-3.2-3B-Instruct (secondary), under identical frozen code, corpus and thresholds.	Tests whether the defense generalises across architectures, and validates the design: L0–L3 inspect input only, and returned numerically identical ASR on both models at every ablation step. Every cross-model difference sits in L4 alone , exactly as the architecture predicts.
Runs and statistics	A single run. No variance estimate, no significance test.	5 seeds × 2 models . Internal ASR 9.77 % ± 1.08 % (Mistral) and 10.33 % ± 0.88 % (Llama); Wilson 95 % intervals and McNemar's test reported throughout ($\chi^2 = 901.0$, $p < 0.001$).	A single run cannot support a claim of stability. The two models' intervals overlap, so the difference between them is not statistically meaningful — which is itself the finding worth reporting.

Changes made to the evaluation methodology, the models tested, and the reported figures, following the reviewers' report

Change	Before	After	Reason
External validation	None. All attacks and all benign queries were written by the author.	BIPIA (Microsoft, KDD '25): 986 attacks and 333 benign documents, plus a fresh never-tuned holdout and 300 real human-written queries drawn from Natural Questions, SciFact and FiQA.	An independent benchmark the author had no role in designing is the only structural answer to co-design bias. End-to-end, 84.08 % of BIPIA attacks were neutralised on Mistral-7B.
Defense code state	Modified during the evaluation period.	Frozen since a single commit; no defense module has changed since. All reported figures come from one identical code state.	Without a freeze, results from different dates are not comparable to each other and cannot be reproduced from the repository as submitted.
FINDINGS RECORDED BUT DELIBERATELY NOT ACTED ON			
External false positives	1.80 % (6 of 333), cause unknown.	Fully root-caused: all six are dense markdown tables. The special-character-ratio rule counts table punctuation as obfuscation evidence, raising LO risk to HIGH, which in turn lowers L3's own threshold. A validated fix takes it to 0.30 % at a cost of one missed attack in 1,001 — and was not applied .	Applying it would invalidate every reported result and require re-running the entire evaluation. It is documented as a diagnosed, quantified limitation with a known remedy rather than folded in silently after the fact.
Reliability of the Llama compliance figure	—	Reported (23.73 %) with its instrument limits stated: 28.7 % of its cases are ambiguous and the figure moves 17.75 points across the margin sweep, against 1.9 % and 2.44 points for Mistral.	The classifier discriminates poorly on this model's responses. Stating that is a limitation of the <i>measurement</i> , not a security claim about the model, and it prevents the weaker number from being read as equivalent to the stronger one.

Every figure above is reproduced from a run of the frozen evaluation code; none is estimated. Several numbers are worse than those originally submitted — the internal false-positive rate moved from 0.00 % to 0.12 % ± 0.16 %, and *semantic camouflage* attack success from 39.9 % to 59.1 % — because the earlier evaluation was structurally unable to expose the errors it contained. Full change history, root cause and verification for each item: revision log, parts 1–4.